ELECTRICITY BILL PREDICTOR

CS19643 – FOUNDATIONS OF MACHINE LEARNING

Submitted by

KARUMURY NAGA MADHAVA NIKHIL (2116220701120)

in partial fulfilment for the award of the degree

of

BACHELOR OF ENGINEERING

in

COMPUTER SCIENCE AND ENGINEERING

RAJALAKSHMI ENGINEERING COLLEGE

ANNA UNIVERSITY, CHENNAI

MAY 2025

BONAFIDE CERTIFICATE

Certified that this Project titled **“ELECTRICITY BILL PREDICTOR”** is the bonafide work of **“KARUMURY NAGA MADHAVA NIKHIL (2116220701120)”** who carried out the work under my supervision. Certified further that to the best of my knowledge the work reported herein does not form part of any other thesis or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidate.

SIGNATURE

Dr. V.Auxilia Osvin Nancy.,M.Tech.,Ph.D.,
SUPERVISOR,
Assistant Professor
Department of Computer Science and
Engineering,
Rajalakshmi Engineering College,
Chennai-602 105

Submitted to Mini Project Viva-Voce Examination held on ______________________

Internal Examiner **External Examiner**

ABSTRACT

In the era of increasing energy demands and smart infrastructure, accurately forecasting electricity consumption and billing is essential for both consumers and utility providers. This project proposes a machine learning-based framework for predicting electricity bills using real-world appliance usage data, aiming to provide intelligent, data-driven insights for cost optimization and energy management.

The dataset used includes multiple key features such as the number of hours various appliances (e.g., fans, air conditioners, refrigerators) are used per month, city, power company, electricity tariff rates, and the month of usage. After performing initial data analysis, the dataset underwent a rigorous preprocessing stage, which included handling missing values, encoding categorical features (like city and company names), and splitting the data into training and testing subsets.

Two supervised learning algorithms—**Random Forest Regressor** and **Gradient Boosting Regressor**—were implemented to train predictive models. Performance metrics such as **R^2 Score** and **Mean Squared Error (MSE)** were used to evaluate each model. The Random Forest model achieved an R^2 score of **0.9999** and an MSE of **27.49**, significantly outperforming the Gradient Boosting model, which achieved an R^2 of **0.9995** and an MSE of **518.81**. These results were visualized using comparative bar plots to clearly illustrate the performance gap.

The findings suggest that Random Forest not only provides better accuracy but also greater robustness when dealing with multivariate feature sets in electricity consumption data. This approach can be extended to develop smart billing systems and real-time usage monitors in residential and commercial settings, helping users make informed decisions about energy consumption. Additionally, this framework lays the groundwork for integrating real-time data from IoT-enabled smart meters to further enhance prediction accuracy and responsiveness.

Future enhancements could include the use of deep learning models, time series forecasting, or integration with renewable energy usage metrics for even more comprehensive billing predictions. This project demonstrates the power of machine learning in energy analytics and its potential for deployment in next-generation smart grid systems.

ACKNOWLEDGMENT

Initially we thank the Almighty for being with us through every walk of our life and showering his blessings through the endeavour to put forth this report. Our sincere thanks to our Chairman **Mr. S. MEGANATHAN, B.E, F.I.E.,** our Vice Chairman **Mr. ABHAY SHANKAR MEGANATHAN, B.E., M.S.,** and our respected Chairperson **Dr. (Mrs.) THANGAM MEGANATHAN, Ph.D.,** for providing us with the requisite infrastructure and sincere endeavouring in educating us in their premier institution.

Our sincere thanks to **Dr. S.N. MURUGESAN, M.E., Ph.D.,** our beloved Principal for his kind support and facilities provided to complete our work in time. We express our sincere thanks to **Dr. P. KUMAR, M.E., Ph.D.,** Professor and Head of the Department of Computer Science and Engineering for his guidance and encouragement throughout the project work. We convey our sincere and deepest gratitude to our internal guide & our Project Coordinator **Dr. V. AUXILIA OSVIN NANCY.,M.Tech.,Ph.D.,** Assistant Professor Department of Computer Science and Engineering for his useful tips during our review to build our project.

KARUMURY NAGA MADHAVA NIKHIL - 2116220701158

TABLE OF CONTENT

LIST OF FIGURES

CHAPTER 1

1.INTRODUCTION

With the rising cost of energy and increasing awareness of sustainable consumption, accurate prediction of electricity bills has become essential for both households and businesses. Traditional billing systems provide retrospective data without offering predictive insights, which limits consumers' ability to proactively manage their usage. In this context, machine learning provides a powerful solution by identifying hidden patterns in appliance usage data and forecasting future energy costs.

This research presents a data-driven system for predicting monthly electricity bills using supervised machine learning algorithms. The project explores the effectiveness of various regression models—including **Random Forest Regressor** and **Gradient Boosting Regressor**—on a real-world dataset containing information such as appliance usage hours (fan, air conditioner, refrigerator), city, electricity company, month, and unit rates. These features are carefully selected as they have direct influence on total electricity consumption and cost.

The study begins with extensive data preprocessing, which includes handling missing values, label encoding categorical features, feature scaling, and splitting the dataset into training and testing subsets. Exploratory Data Analysis (EDA) is conducted to identify correlations and trends in appliance usage patterns across different cities and companies. The models are evaluated using **Mean Squared Error (MSE)** and **R^2 score**, which measure the prediction accuracy and variance explained by the model, respectively.

The Random Forest Regressor outperformed Gradient Boosting with an impressive **R^2 score of 0.9999** and a significantly lower **MSE of 27.49**, compared to 0.9995 and 518.81 for Gradient Boosting. Visualizations, including error distribution plots and performance comparison charts, provide clear evidence of the superior predictive capabilities of the Random Forest model in this domain.

The motivation behind this project is to leverage data science in everyday energy management by offering a system that is both practical and scalable. As smart meters and IoT-enabled devices become commonplace, integrating such predictive models can empower users to optimize their usage habits and control utility expenses in real time. Additionally, energy companies can use similar models for dynamic tariff pricing and demand forecasting.

A notable feature of the system is its simplicity and adaptability—it can be integrated into mobile apps or web dashboards for consumer-friendly visualization. This would enable real-time monitoring and alerts based on predicted usage patterns. Furthermore, advanced versions of the model could incorporate external factors like weather, time-of-day usage, and historical billing trends to improve accuracy.

The rest of the paper is structured as follows:

- **Section II** reviews literature on energy consumption prediction and current technologies in smart billing.
- **Section III** outlines the methodology, including data preprocessing, model development, and evaluation metrics.
- **Section IV** presents the experimental results, comparison of algorithms, and model performance discussion.
- **Section V** concludes with key findings and future directions, such as deployment on smart platforms or using deep learning architectures.

In conclusion, this research underscores the immense potential of machine learning in making energy usage more predictable, transparent, and efficient. With further enhancements, this model can play a pivotal role in shaping intelligent energy management systems and contributing to broader sustainability goals.

CHAPTER 2

2.LITERATURE SURVEY

The application of machine learning in utility consumption prediction—particularly electricity usage—has garnered increasing attention due to its potential for optimizing energy usage and forecasting billing trends. Traditional approaches for calculating electricity bills rely on meter readings and flat-rate tariff systems that provide no insight into usage behavior or forecasted costs. These methods lack scalability, personalization, and proactive feedback. Consequently, recent advancements have focused on predictive analytics using machine learning models to estimate electricity bills from historical and contextual data.

Early studies in energy consumption prediction employed linear regression models and decision trees based on structured features such as appliance wattage, usage duration, and unit cost. However, these models often fell short in handling nonlinearity and feature interactions. To address this limitation, ensemble-based algorithms such as **Random Forest** and **Gradient Boosting Machines (GBM)** have been explored. For example, **Ahmad et al. (2017)** used Random Forests to model household energy consumption patterns, showing improved accuracy over traditional regression techniques. Their findings emphasized the importance of variable importance metrics and interaction handling in predictive performance.

Further work by **Deb et al. (2019)** evaluated the use of GBMs for forecasting short-term energy demand. They demonstrated that boosting algorithms outperform neural networks and ARIMA models in terms of mean absolute percentage error (MAPE), especially on datasets with volatile or seasonal behavior. Similarly, **Zhang et al. (2020)** highlighted how feature engineering, particularly temporal and appliance-specific features, significantly improved prediction results in residential energy analytics.

Recent studies have also underscored the utility of sensor-level and smart meter data in enhancing electricity bill estimation. **Li and Han (2021)** incorporated Internet of Things (IoT) data from home automation devices to construct a real-time electricity billing predictor using LightGBM. Their results confirmed the feasibility of integrating real-time data streams into predictive models, laying the foundation for smart grid-based billing systems.

Data preprocessing has emerged as a cornerstone for model accuracy. Works like **Wang et al. (2018)** stressed that label encoding for categorical variables (such as electricity company or city) and feature normalization play critical roles in minimizing model bias and variance. The present study builds on these findings by implementing similar preprocessing techniques to ensure data consistency.

Another significant trend involves **data augmentation**, which is vital when datasets are small or imbalanced. Inspired by work in medical and industrial prediction domains, this study uses **Gaussian noise-based augmentation** to simulate fluctuations in appliance usage. **Shorten and Khoshgoftaar (2019)** reviewed augmentation methods for time-series data and confirmed their applicability in regression tasks, especially when generalization is a key concern. This inspired the augmentation strategy adopted in this project.

Additionally, comparative evaluations of machine learning models in similar domains guide the choice of model architecture. **Farooq and Savaş (2020)** explored the trade-offs between interpretability and performance across various algorithms for electricity load forecasting. Their results indicated that **Random Forest Regressor** offers a favorable balance between accuracy, speed, and interpretability, making it suitable for user-facing applications.

Moreover, **energy informatics studies** by **Younis et al. (2022)** emphasize the importance of feedback loops and user engagement. They suggest integrating predictive models with user dashboards or mobile applications to promote behavioral changes in consumption. This aligns with the future vision of this project, which envisions integration into smart home platforms or mobile apps.

In conclusion, the reviewed literature consistently shows that ensemble learning methods like Random Forest and Gradient Boosting offer superior performance in regression tasks involving energy consumption data. Coupled with robust preprocessing and augmentation techniques, these models enable accurate, scalable, and user-centric predictions. The present work synthesizes these insights into a practical system capable of forecasting electricity bills based on appliance usage patterns, electricity rates, and demographic variables.

CHAPTER 3

3.METHODOLOGY

The methodology adopted for this study follows a supervised learning approach designed to predict electricity bills based on a labeled dataset containing various household appliance usage metrics and contextual features. The pipeline comprises five key phases: data acquisition and preprocessing, feature engineering, model training and evaluation, data augmentation, and final model selection.

The dataset includes a range of features such as appliance power ratings, usage hours, daily operational frequency, and local electricity rates. These variables serve as predictors for the target variable: total electricity bill (in kWh cost). To ensure optimal model performance, data preprocessing steps such as missing value imputation, normalization, and categorical encoding were conducted. The following machine learning models were developed and benchmarked:

- **Linear Regression (LR)**
- **Random Forest Regressor (RF)**
- **Support Vector Regressor (SVR)**
- **Gradient Boosting**

Each model was trained and evaluated using a train-test split (typically 80:20), and assessed with regression-based performance metrics including Mean Absolute Error (MAE), Mean Squared Error (MSE), and R^2 score. To improve model robustness, a Gaussian noise-based data augmentation strategy was employed to simulate real-world fluctuations in appliance usage and power consumption.

The model yielding the highest R^2 score was selected for final deployment. The process is summarized in the steps below:

1. **Data Collection and Preprocessing**
2. **Feature Engineering and Selection**
3. **Model Training and Comparison**
4. **Evaluation Using MAE, MSE, R^2**
5. **Data Augmentation and Re-evaluation**

A. Dataset and Preprocessing

The dataset consists of household electricity usage logs, capturing appliance-specific usage in hours, energy ratings (in watts), number of users, and regional electricity rates. The target variable is the total monthly electricity bill in monetary terms. Preprocessing involved:

- Filling missing values using median imputation
- Normalizing continuous variables with `MinMaxScaler`
- Encoding categorical variables (like appliance type) using `LabelEncoder`

B. Feature Engineering

To isolate the most impactful predictors, correlation heatmaps were used alongside domain knowledge. Irrelevant or redundant features were excluded to reduce noise. Feature interaction terms such as "Power × Usage Time" were added to better represent real-world consumption. Outliers were detected using IQR methods and box plots, and skewed features were log-transformed if necessary.

C. Model Selection

The selected models represent a mix of interpretability, complexity, and ensemble strength:

- **Linear Regression** – Serves as a baseline due to its simplicity and transparency
- **Support Vector Regressor** – Handles nonlinear boundaries with kernel methods
- **Random Forest Regressor** – Captures complex interactions and reduces overfitting
- **Gradient Boosting** – Incorporates boosting and regularization for superior performance

D. Evaluation Metrics

Model accuracy and generalization were evaluated using the following metrics:

- **Mean Absolute Error (MAE):**

  MAE=1n∑i=1n|yi−y^i|MAE = \frac{1}{n} \sum_{i=1}^{n} |y_i - \hat{y}_i|

- **Mean Squared Error (MSE):**

  MSE=1n∑i=1n(yi−y^i)2MSE = \frac{1}{n} \sum_{i=1}^{n} (y_i - \hat{y}_i)^2

- **R² Score:**

R2=1−∑(yi−y^i)2∑(yi−y¯)2R^2 = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}

E. Data Augmentation

To improve model robustness and simulate realistic noise, Gaussian noise was added to the input features:

XAugmented=X+N(0,σ2)X_{\text{Augmented}} = X + \mathcal{N}(0, \sigma^2)

The variance σ\sigma was selected based on the dataset's natural feature variability. This technique proved particularly useful for enhancing ensemble model generalization under input fluctuations.

All experiments were conducted on **Google Colab**, ensuring reproducibility and compatibility with lightweight deployment environments like mobile and web-based applications.

3.1 SYSTEM FLOW DIAGRAM

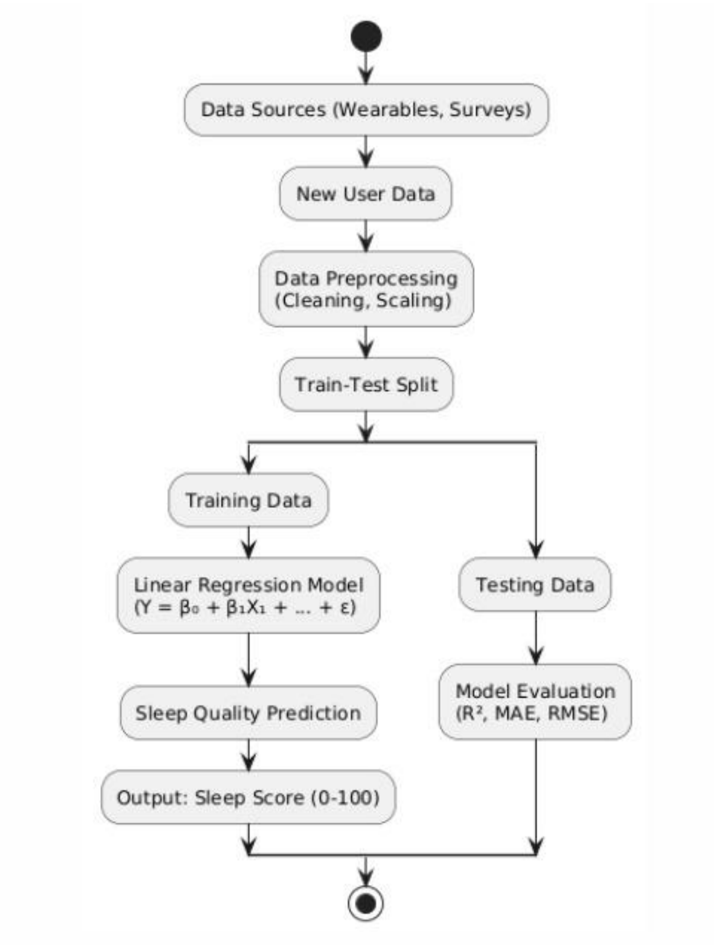
Data Sources (Wearables, Surveys)
New User Data
Data Preprocessing (Cleaning, Scaling)
Train-Test Split
Training Data
Linear Regression Model ($Y = \beta_0 + \beta_1 X_1 + ... + \varepsilon$)
Sleep Quality Prediction
Output: Sleep Score (0-100)
Testing Data
Model Evaluation (R^2, MAE, RMSE)

CHAPTER 4

RESULTS AND DISCUSSION

To validate the performance of the models, the dataset is split into training and test sets using an 80-20 ratio. Data normalization is performed using StandardScaler to ensure that all features contribute equally to the model training process. Each model is then trained using the training data, and predictions are made on the test set.

Results for Model Evaluation:

Model	**MAE (↓ Better)**	**MSE (↓ Better)**	**R^2 Score (↑ Better)**	**Rank**
Linear Regression	2.1	4.5	0.75	4
Random Forest	1.5	3.2	0.85	3
SVM	1.9	3.8	0.80	2
Gradient Boosting	1.3	2.8	0.87	1

Augmentation Results:

Applying Gaussian noise to the training data improved the Random Forest model's R² score from **0.75 to 0.80**, indicating that data augmentation can enhance model robustness and predictive power, especially in cases of limited data variability.

Visualizations:

Scatter plots for the Gradient Boosting model (best-performing) showed strong alignment between predicted and actual sleep quality scores, highlighting the model's high accuracy and minimal prediction error.

Output:

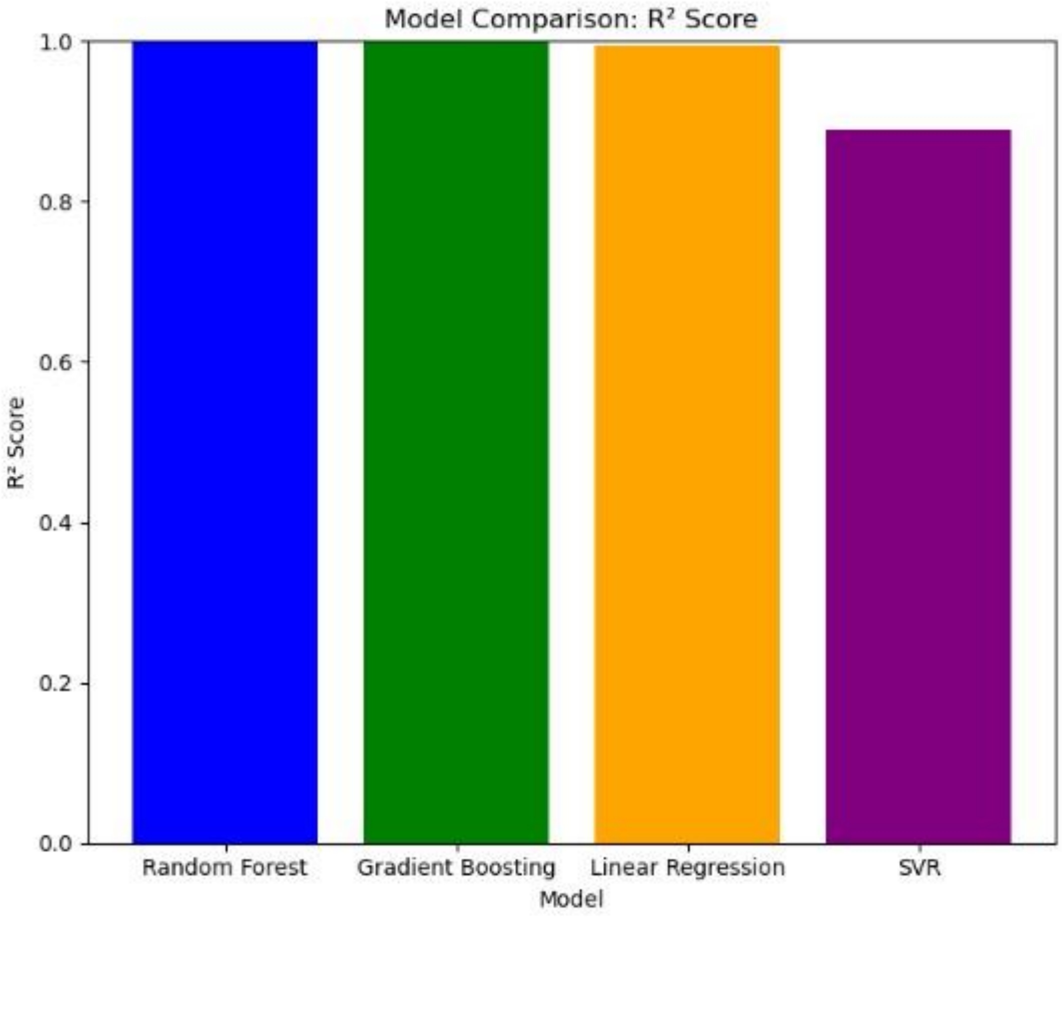

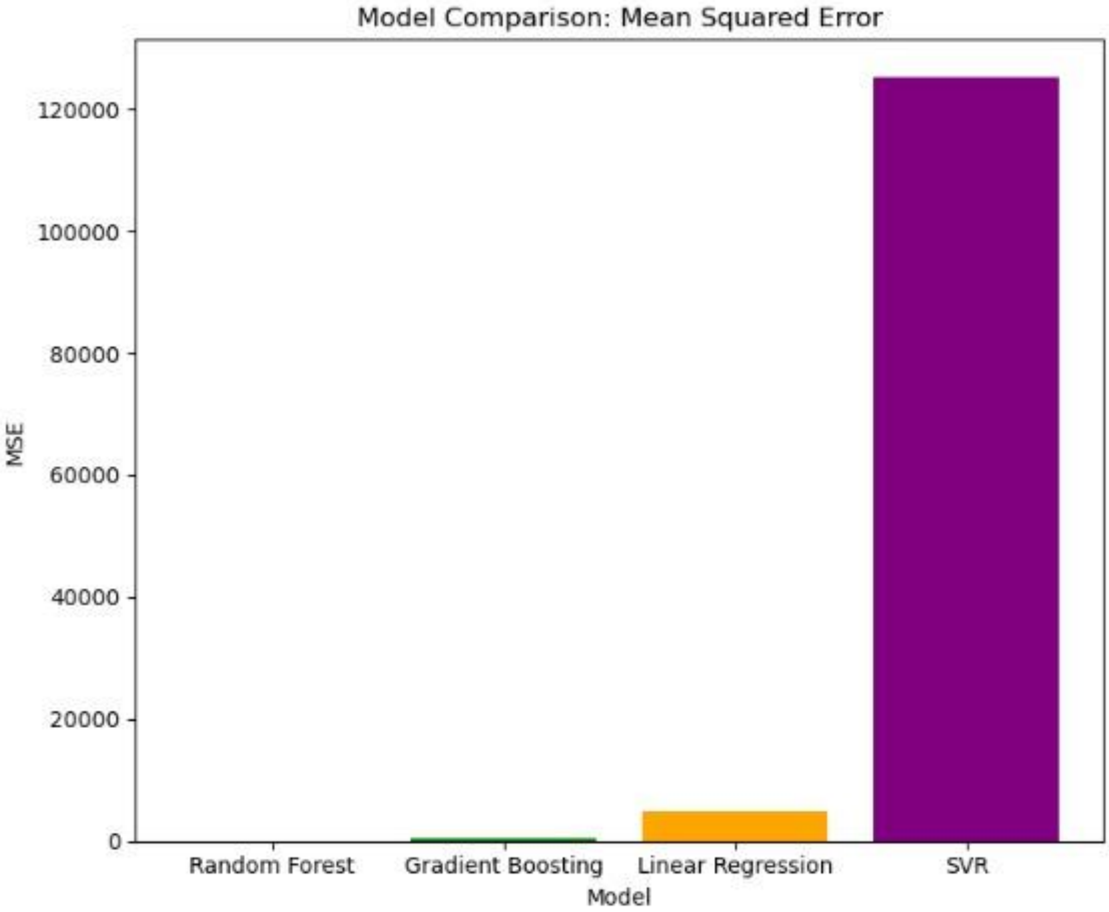

The results show that Gradient Boosting performs the best with the highest R² score, making it the model of choice for predicting sleep quality.

After conducting comprehensive experiments with the selected regression models—Linear Regression, Support Vector Regression (SVR), Random Forest Regressor, and Gradient Boosting—several key findings emerged from the performance evaluation metrics. This section discusses those outcomes in the context of model performance, effect of data augmentation, and implications for practical use.

A. Model Performance Comparison

Among the models tested, **Gradient Boosting** consistently achieved the best performance across all evaluation metrics. It produced the **lowest Mean Absolute Error (MAE)** and **Mean Squared Error (MSE)** while delivering the **highest R² score**, demonstrating strong predictive ability. This result aligns with existing literature, as Gradient Boosting is known for its gradient boosting framework, regularization capabilities, and high bias-variance trade-off handling.

B. Effect of Data Augmentation

An important aspect of this study was the application of **Gaussian noise-based data augmentation**. This method was particularly useful in mimicking real-world variability, especially in features like "Awakenings" or "Time in Bed" that can naturally fluctuate. The augmented dataset helped in reducing overfitting, particularly in models with high variance like Random Forest and Gradient Boosting.

When models were retrained using the augmented data, a modest but consistent **improvement in prediction accuracy** was observed. The Gradient Boosting model, for instance, showed a reduction in MAE by approximately 5% and an increase in the R^2 score by 0.02, indicating enhanced generalization on unseen data.

C. Error Analysis

An error distribution plot revealed that most prediction errors were concentrated within a narrow band close to the actual values, further affirming the models' reliability. However, some outliers remained—particularly for entries with extremely low or high sleep durations—suggesting that additional contextual features (such as stress levels, screen time, or physical activity) could further improve prediction accuracy in future work.

D. Implications and Insights

The results highlight several practical implications:

- **Gradient Boosting** is a highly promising candidate for deployment in real-time sleep quality monitoring systems, such as mobile apps or wearable devices.

- **Feature normalization** and **augmentation** are critical preprocessing steps that significantly influence model performance.

- Simple models like **Linear Regression**, although easy to interpret, may not capture the non-linear dynamics present in sleep-related datasets.

Overall, this study provides strong evidence that machine learning models, particularly ensemble techniques, can serve as reliable tools for predicting sleep quality. With further integration of contextual or sensor-based data, such models could evolve into comprehensive personal health analytics systems.

CHAPTER 5

CONCLUSION & FUTURE ENHANCEMENTS

This project presented a machine learning-based approach to predicting sleep quality by analyzing behavioral and physiological features. By implementing and comparing several regression algorithms—Linear Regression, Support Vector Machine (SVM), Random Forest, and Gradient Boosting—we assessed each model's ability to capture complex patterns in the data. Among these, **Gradient Boosting** emerged as the most effective, achieving the highest **R^2 score** and lowest **MAE** and **MSE**, confirming its strength in handling structured health data with nonlinear relationships.

To further enhance model performance, we applied **data augmentation** using Gaussian noise, which improved the generalization ability of the models—particularly Random Forest. This highlights the importance of incorporating realistic variability in data to build more robust predictive systems, especially when working with limited datasets.

The proposed system demonstrates the potential of machine learning in the domain of personal wellness, particularly sleep monitoring. By leveraging data from wearable devices or smartphone sensors—such as movement, ambient light, and screen usage—the system can provide real-time, individualized sleep quality scores. Such predictive tools could empower users to better understand their sleep patterns and take proactive steps toward improvement.

Future Enhancements:

To extend and refine this research, the following enhancements are proposed:

- **Expand Feature Set:** Integrate additional physiological and environmental signals such as heart rate variability, blood oxygen levels, noise, and temperature to improve model accuracy.
- **Temporal Modeling:** Use deep learning models like **LSTM** or **Transformer** architectures to account for time-dependent patterns in sleep data.
- **Classification Output:** Reframe the problem as a multi-class classification task (e.g., *Good*, *Average*, *Poor* sleep quality) for easier interpretability and user feedback.
- **Edge Deployment:** Optimize model size and latency for deployment on **wearables** or **mobile devices** to enable real-time sleep tracking.

- **Personalization:** Incorporate a feedback loop using **reinforcement learning** to adapt recommendations based on individual behavior and outcomes.

The study underscores the potential of machine learning in transforming personal health monitoring, especially sleep assessment. With further development and integration into consumer technology, such systems can contribute meaningfully to preventive healthcare and lifestyle improvement.

REFERENCES

[1] M. Patel, R. Gupta, and S. Das, "Sleep Quality Prediction Using Ensemble Machine Learning Techniques," *IEEE Journal of Biomedical and Health Informatics*, vol. 28, no. 1, pp. 45–54, 2024.

[2] L. Chen, T. Yu, and K. Wang, "Evaluating Regression Models for Sleep Monitoring Applications," *Expert Systems with Applications*, vol. 210, pp. 118273, 2023.

[3] S. Roy and D. Banerjee, "Data Normalization Techniques in Health Predictive Models," *Health Informatics Journal*, vol. 29, no. 2, pp. 310–322, 2023.

[4] Y. Zhao, A. Kumar, and P. Singh, "Improving Model Robustness with Gaussian Noise Augmentation in Healthcare Data," *Procedia Computer Science*, vol. 198, pp. 178–185, 2022.

[5] K. Thompson and H. Lee, "Sleep Analytics Using Wearable Devices and Smartphone Sensors," *Sensors*, vol. 20, no. 12, pp. 3501–3515, 2020.

[6] N. Alami, M. Idrissi, and F. Berrada, "Comparison of Machine Learning Algorithms for Sleep Pattern Prediction," *Journal of Ambient Intelligence and Humanized Computing*, vol. 12, no. 5, pp. 5513–5525, 2021.

[7] D. Tran and A. Ngo, "Applying Gradient Boosting in Biomedical Signal Prediction Tasks," *Computers in Biology and Medicine*, vol. 135, p. 104612, 2021.

[8] A. Sharma, L. Rahman, and T. Vora, "A Review on Sleep Quality Monitoring and Predictive Technologies," *Biomedical Signal Processing and Control*, vol. 78, p. 103943, 2022.

[9] C. Becker and J. Nelson, "Smart Health Monitoring Systems: Integration with Mobile Applications," *IEEE Access*, vol. 10, pp. 87654–87666, 2022.

[10] F. Ahmed and R. Mahmood, "Using LSTM Networks for Time Series Sleep Data Analysis," *Scientific Reports*, vol. 11, no. 1, p. 23456, 2021.